Package 'rapidsplithalf'

January 27, 2025

Type Package

Title A Fast Permutation-Based Split-Half Reliability Algorithm

Version 0.3

Date 2024-08-14

Description Accurately estimates the reliability of cognitive tasks using a fast and flexible permutation-based split-half reliability algorithm that supports stratified splitting while maintaining equal split sizes. See Kahveci, Bathke, and Blechert (2022) <doi:10.31234/osf.io/ta59r> for details.

License GPL (>= 2)

BugReports https://github.com/Spiritspeak/rapidsplit/issues/

Depends R(>= 4.0)

Imports Rcpp (>= 1.0.5), doParallel, foreach

Suggests knitr, rmarkdown

LinkingTo Rcpp

RoxygenNote 7.3.2

Encoding UTF-8

VignetteBuilder knitr

Contents

1

---

bootstrapWeights	*Bootstrap Weights*

Description

Create a matrix of bootstrap samples expressed as frequency weights

Usage

```
bootstrapWeights(size, times)
```

Arguments

`size`	Number of values to bootstrap
`times`	Number of bootstraps

Value

A matrix with bootstrap samples expressed as frequency weights. Each column represents a single bootstrap iteration and each row represents a case.

Examples

```
# Rapidly compute a bootstrapped median to obtain its standard error
myweights<-bootstrapWeights(size=50, times=100)
meds<-mediansByWeight(x=rnorm(50),weights=myweights)
# SE
sd(meds)
```

---

colAggregators	*Fast matrix column aggregators*

Description

Fast matrix column aggregators

Usage

```
colMedians(x)

colProds(x)

colSds(x)

colMediansMasked(x, mask)

colMeansMasked(x, mask)

colSdsMasked(x, mask)
```

Arguments

x	A numeric matrix to compute column aggregates of.
mask	A logical matrix determining which data points to include in the column-wise aggregations.

Value

A numeric vector representing values aggregated by column.

Author(s)

Sercan Kahveci

See Also

colMeans, mediansByMask, maskAggregators

Examples

```
x <- cbind(x1 = 3, x2 = c(4:1, 2:5))
colMedians(x)

colProds(x)

colSds(x)

mask<-cbind(rep(c(TRUE,FALSE),4),
            rep(c(TRUE,FALSE),each=4))
colMediansMasked(x,mask)

colMeansMasked(x,mask)

colSdsMasked(x,mask)
```

---

corByColumns	*Correlate two matrices by column*

Description

Correlate each column of 1 matrix with the same column in another matrix

Usage

```
corByColumns(x, y)

corByColumns_mask(x, y, mask)

corStatsByColumns(x, y)
```

Arguments

`x, y`	Matrices whose values to correlate by column.
`mask`	Logical matrix marking which data points to include.

Details

The primary use for these functions is to rapidly compute the correlations between two sets of split-half scores stored in matrix columns.

`corStatsByColumns` produces the mean correlation of all column-pairs using the formula `mean(covariances) / sqrt(mean(col1variance) * mean(col2variance))`

This method is more accurate than [cormean()](#) and was suggested by prof. John Christie of Dalhousie University.

Value

`corByColumns()` and `corByColumns_mask()` return a numeric vector of correlations of each pair of columns.

`corStatsByColumns()` returns a list with named items:

- cormean: the aggregated correlation coefficient of all column pairs (see Details)
- allcors: the correlations of each column pair
- xvar: the column variances of matrix x
- yvar: the column variances of matrix y
- covar: the covariances of each column pair

Author(s)

Sercan Kahveci

Examples

```
m1<-matrix((1:9)+rnorm(9),ncol=3)
m2<-matrix((9:1)+rnorm(9),ncol=3)
corByColumns(m1,m2)

mask<-1-diag(3)
corByColumns_mask(m1,m2,mask)

corStatsByColumns(m1,m2)
```

---

cormean	*Compute a minimally biased average of correlation values*

Description

This function computes a minimally biased average of correlation values. This is needed because simple averaging of correlations is negatively biased, and the often used z-transformation method of averaging correlations is positively biased. The algorithm was developed by Olkin & Pratt (1958).

Usage

```
cormean(
  r,
  n,
  weights = c("none", "n", "df"),
  type = c("OP5", "OP2", "OPK"),
  na.rm = FALSE,
  incl.trans = FALSE
)
```

Arguments

r	A vector containing correlation values/
n	A single value or vector containing sample sizes/
weights	Character. How should the correlations be weighted? none leads to no weighting, n weights by sample size, df weights by sample size minus one.
type	Character. Determines which averaging algorithm to use, with "OP5" usually being the most accurate.
na.rm	Logical. Should missing values be removed?
incl.trans	Logical. Should the transformed correlations be included?

Value

An average correlation.

References

Olkin, I., & Pratt, J. (1958). Unbiased estimation of certain correlation coefficients. The Annals of Mathematical Statistics, 29. https://doi.org/10.1214/aoms/1177706717

Shieh, G. (2010). Estimation of the simple correlation coefficient. Behavior Research Methods, 42(4), 906-917. https://doi.org/10.3758/BRM.42.4.906

Examples

```
cormean(c(0,.3,.5),c(30,30,60))
```

---

correlation-tools *Miscellaneous correlation tools*

Description

Helper functions to compute important statistics from correlation coefficients.

Usage

```
r2z(r)

z2r(z)

r2t(r, n)

t2r(t, n)

r2p(r, n)

rconfint(r, n, alpha = 0.05)

compcorr(r1, r2, n1, n2)

## S3 method for class 'compcorr'
print(x, ...)
```

Arguments

r, r1, r2	Correlation values.
z	Z-scores.
n, n1, n2	Sample sizes.
t	t-scores.
alpha	The significance level to use.
x	A compcorr object to print.
...	Ignored.

Value

For r2z(), z2r, r2t, t2r, and r2p, a numeric vector with the requested transformation applied. For rconfint(), a numeric vector with two values representing the lower and upper confidence intervals of the correlation coefficient. For compcorr(), a compcorr object containing a z and p value for the requested comparison, which can be printed with print.compcorr().

Functions

- r2z(): Converts correlation coefficients to z-scores.
- z2r(): Converts z-scores to correlation coefficients.
- r2t(): Converts correlation coefficients to t-scores.
- t2r(): Converts t-scores to correlation coefficients.
- r2p(): Computes the two-sided p-value for a given correlation.
- rconfint(): Computes confidence intervals for one or multiple correlation coefficients.
- compcorr(): Computes the significance of the difference between two correlation coefficients.
- print(compcorr): Computes the significance of the difference between two correlation coefficients.

Author(s)

Sercan Kahveci

See Also

cormean

Examples

```
z <- r2z(.5)
r <- z2r(z)
t<-r2t(r,30)
r<-t2r(t,30)
r2p(r,30)
print(rconfint(r,30))
print(compcorr(.5,.7,20,20))
```

---

excludeOutliersByMask   *Exclude SD-based outliers*

Description

Different masks (columns of a logical matrix) are applied to the same input vector, and outliers in each resulting subvector are marked with FALSE in the mask.

Usage

```
excludeOutliersByMask(x, mask, sdlim = 3)
```

Arguments

x	Vector to exclude outliers from.
mask	A logical matrix determining which data points to include and which not to.
sdlim	Standard deviation limit to apply; values beyond are classified as outliers and masked.

Value

An updated mask.

Examples

```
x<-rnorm(50)
x[1]<-100
x[2]<-50
mask<-matrix(TRUE,ncol=3,nrow=50)
mask[1,2]<-FALSE
mask[2,3]<-FALSE
excludeOutliersByMask(x,mask)
```

foodAAT	*Approach-Avoidance Task examining approach bias to different foods*

Description

This data originates from an approach-avoidance task examining approach bias towards food. Participants responded to the stimulus category (food or object) by pulling or pushing a joystick. Instructions were flipped from one block to the next.

Usage

```
data(foodAAT)
```

Format

An object of class `"data.frame"`.

Details

- subjectid: Participant ID.
- stimid: Stimulus ID.
- is_pull: Whether the trial required an approach response (1) or an avoid response (0).
- is_target: Whether the trial featured a food stimulus (1) or an object stimulus (0).
- error: Whether the response was incorrect (1) or correct (0).
- RT: The response initiation time.
- FullRT: The time from stimulus onset to response completion.
- trialnum: The trial number.
- blocknum: The block number.
- palatability: The participant's palatability rating for the stimulus (foods only).
- valence: The participant's valence rating for the stimulus.
- FCQS_2_craving: The participant's FCQS state food craving score at time of testing.
- FCQS_2_hunger: The participant's FCQS state hunger score at time of testing.

Source

References

Lender, A., Meule, A., Rinck, M., Brockmeyer, T., & Blechert, J. (2018). Measurement of food-related approach–avoidance biases: Larger biases when food stimuli are task relevant. Appetite, 125, 42–47. doi:10.1016/j.appet.2018.01.032

---

generateSplits   *A balanced split-half generator*

Description

Generates split-half indices that can be stratified by multiple subgroup variables while guaranteeing near-equal numbers of trials in both halves.

Usage

```
generateSplits(data, subsetvars, stratvars = NULL, splits, verbose = TRUE)
```

Arguments

data	A dataset to generate split-halves from.
subsetvars	Variables identifying subgroups that must be individually split into equally sized halves, e.g. participant number and experimental condition.
stratvars	Variables identifying subgroups that are nested within the subsetvars, and must be split as fairly as possible, while preserving the equal size of the two halves of each subset identified by the subsetvars, e.g. stimulus ID.
splits	How many splits to generate.
verbose	Display progress bar?

Value

A logical `matrix` in which each row represents a row of the input dataset, and each column represents a single split.

Examples

```
data(foodAAT)
mysplits<-generateSplits(data=foodAAT,
                         subsetvars=c("subjectid","is_pull","is_target"),
                         stratvars="stimid",
                         splits=1)
half1<-foodAAT[ mysplits[,1],]
half2<-foodAAT[!mysplits[,1],]
```

---

maskAggregators   *Multi-mask/weight based aggregators*

Description

Methods to aggregate the same vector with different masks or frequency weights. Useful for fast bootstrapping or split-half scoring. A single aggregate value of x is computed for each column of the mask or weight matrix.

Usage

```
mediansByMask(x, mask)

meansByMask(x, mask)

sdsByMask(x, mask)

mediansByWeight(x, weights)

meansByWeight(x, weights)

sdsByWeight(x, weights)
```

Arguments

x	A vector to aggregate over with different masks or weights.
mask	Logical matrix where each column represents a separate vector of masks to aggregate x with. Only values marked `TRUE` are included in the aggregation.
weights	Integer matrix where each column represents frequency weights to weight the aggregation by.

Value

a vector with each value representing an aggregate of the same single input vector but with different masks or frequency weights applied.

Author(s)

Sercan Kahveci

See Also

colMedians, colAggregators, generateSplits

Examples

```
# Demonstration of mediansByMask()
x<-1:6
mask<-rbind(c(TRUE,FALSE,FALSE),
            c(TRUE,FALSE,FALSE),
            c(FALSE,TRUE,FALSE),
            c(FALSE,TRUE,FALSE),
            c(FALSE,FALSE,TRUE),
            c(FALSE,FALSE,TRUE))
mediansByMask(x,mask)

# Compute split-halves for a single
# participant, stratified by stimulus
data(foodAAT)
currdata<-foodAAT[foodAAT$subjectid==3,]
currdata$stratfactor<-
  interaction(currdata$is_pull,
              currdata$is_target,
              currdata$stimid)
```

```
currdata<-currdata[order(currdata$stratfactor),]
groupsizes<-
  rle(as.character(currdata$stratfactor))$lengths
mysplits<-
  stratifiedItersplits(splits=1000,
                       groupsizes=groupsizes)

# Median for half 1
mediansByMask(currdata$RT,mysplits==1)

#How to use meansByMask()
meansByMask(x,mask)
sd(meansByMask(currdata$RT,mysplits==1))

# How to use sdsByMask() to compute
# mask-based D-scores
meansByMask(currdata$RT,mysplits==1) /
  sdsByMask(currdata$RT,mysplits==1)

# Compute the bootstrapped
# standard error of a median
weights<-
  bootstrapWeights(size=nrow(currdata),
                   times=1000)
bootmeds<-mediansByWeight(currdata$RT,weights)
sd(bootmeds) # bootstrapped standard error

# Compute the bootstrapped
# standard error of a mean
bootmeans<-meansByWeight(currdata$RT,weights)
sd(bootmeans) # bootstrapped standard error
# exact standard error for comparison
sd(currdata$RT)/sqrt(length(currdata$RT))

# Use sdsByWeight to compute bootstrapped D-scores
bootsds<-sdsByWeight(currdata$RT,weights)
# bootstrapped standard error of D-score
sd(bootmeans/bootsds)
```

OutlierMaskers   *Exclude SD-based outliers in each matrix column*

Description

Generate or update a mask matrix based on outlyingness of values in each column.

Usage

```
maskOutliers(x, sdlim = 3)

maskOutliersMasked(x, mask, sdlim = 3)
```

Arguments

x	Matrix in which to mark SD-based outliers by column.
sdlim	Standard deviation limit to apply; values beyond are classified as outliers and masked.
mask	A logical matrix determining which data points to include and which not to.

Value

A logical matrix with outliers (and previously masked values) marked as FALSE.

Examples

```
# Generate data with outliers
testmat<-matrix(rnorm(100),ncol=2)
testmat[1,]<-100
testmat[2,]<-50

# Detect outliers
maskOutliers(testmat)

# Generate a mask
testmask<-matrix(TRUE,ncol=2,nrow=50)
testmask[1,1]<-FALSE

# Detect outliers with pre-existing mask
maskOutliersMasked(x=testmat,
                   mask=testmask, sdlim = 3)
```

---

raceIAT	*Implicit Association Task examining implicit bias towards White and Black people*

Description

This data originates from the publicly available implicit association test (IAT) on racial prejudice hosted by Project Implicit. 200 participants were randomly sampled from the full trial-level data available for participants from 2002 to 2022. We included only those IAT blocks relevant to scoring (3,4,6,7) and only individuals with full data.

Usage

```
data(raceIAT)
```

Format

An object of class "data.frame".

Details

- session_id: The session id, proxy for participant number.
- task_name: Subtype of IAT used.
- block_number: IAT block number.
- block_pairing_definition: Stimulus pairing displayed in block.
- block_trial_number: Trial number within block.
- stimulus: Stimulus name.
- required_response: The response required from the participant.
- latency: Participant's response latency.
- error: Whether the response was wrong (TRUE).
- trial_number: Experimentwise trial number.
- stimcat: The stimulus category.
- respcat: Category of the required response.
- blocktype: Either practice block or full IAT block.
- congruent: Whether the block was congruent with anti-black bias (TRUE) or not.
- latency2: Response latencies with those for incorrect responses replaced by the block mean plus a penalty.

Source

[OSF.io repository](#)

References

Xu, K., Nosek, B., & Greenwald, A. G. (2014). Psychology data from the race implicit association test on the project implicit demo website. Journal of open psychology data, 2(1), e3-e3. [doi:10.5334/jopd.ac](#)

---

rapidsplit   *rapidsplit*

Description

A very fast algorithm for computing stratified permutation-based split-half reliability.

Usage

```
rapidsplit(
  data,
  subjvar,
  diffvars = NULL,
  stratvars = NULL,
  subscorevar = NULL,
  aggvar,
  splits = 6000,
  aggfunc = c("means", "medians"),
```

```
  errorhandling = list(type = c("none", "fixedpenalty"), errorvar = NULL, fixedpenalty =
    600, blockvar = NULL),
  standardize = FALSE,
  include.scores = TRUE,
  verbose = TRUE,
  check = TRUE
)

## S3 method for class 'rapidsplit'
print(x, ...)

## S3 method for class 'rapidsplit'
plot(
  x,
  type = c("average", "minimum", "maximum", "random", "all"),
  show.labels = TRUE,
  ...
)

rapidsplit.chunks(
  data,
  subjvar,
  diffvars = NULL,
  stratvars = NULL,
  subscorevar = NULL,
  aggvar,
  splits = 6000,
  aggfunc = c("means", "medians"),
  errorhandling = list(type = c("none", "fixedpenalty"), errorvar = NULL, fixedpenalty =
    600, blockvar = NULL),
  standardize = FALSE,
  include.scores = TRUE,
  verbose = TRUE,
  check = TRUE,
  chunks = 4,
  cluster = NULL
)
```

Arguments

data	Dataset, a `data.frame`.
subjvar	Subject ID variable name, a `character`.
diffvars	Names of variables that determine which conditions need to be subtracted from each other, `character`.
stratvars	Additional variables that the splits should be stratified by; a `character`.
subscorevar	A `character` variable identifying subgroups within a participant's data from which separate scores should be computed. To compute a participant's final score, these subscores will be averaged together. A typical use case is the D-score of the implicit association task.
aggvar	Name of variable whose values to aggregate, a `character`. Examples include reaction times and error rates.

splits	Number of split-halves to average, an `integer`. It is recommended to use around 5000.
aggfunc	The function by which to aggregate the variable defined in `aggvar`; can be `"means"`, `"medians"`, or a custom function (not a function name). This custom function must take a numeric vector and output a single value.
errorhandling	A list with 4 named items, to be used to replace error trials with the block mean of correct responses plus a fixed penalty, as in the IAT D-score. The 4 items are `type` which can be set to `"none"` for no error replacement, or `"fixedpenalty"` to replace error trials as described; `errorvar` requires name of the `logical` variable indicating an incorrect response (as TRUE); `fixedpenalty` indicates how much of a penalty should be added to said block mean; and `blockvar` indicates the name of the block variable.
standardize	Whether to divide by scores by the subject's SD; a `logical`. Regardless of whether error penalization is utilized, this standardization will be based on the unpenalized SD of correct and incorrect trials, as in the IAT D-score.
include.scores	Include all individual split-half scores?
verbose	Display progress bars? Defaults to `TRUE`.
check	Check input for possible problems?
x	`rapidsplit` object to print or plot.
...	Ignored.
type	Character argument indicating what should be plotted. By default, this plots the random split whose correlation is closest to the average. However, this can also plot the random split with the `"minimum"` or `"maximum"` split-half correlation, or any `"random"` split. `"all"` splits can also be plotted together in one figure.
show.labels	Should participant IDs be shown above their points in the scatterplot? Defaults to `TRUE` and is ignored when `type` is `"all"`.
chunks	Number of chunks to divide the splits in, for more memory-efficient computation, and to divide over multiple cores if requested.
cluster	Chunks will be run on separate cores if a cluster is provided, or an `integer` specifying the number of cores. Otherwise, if the value is `NULL`, the chunks are run sequentially.

Details

The order of operations (with optional steps between brackets) is:

- Splitting
- (Replacing error trials within block within split)
- Computing aggregates per condition (per subscore) per person
- Subtracting conditions from each other
- (Dividing the resulting (sub)score by the SD of the data used to compute that (sub)score)
- (Averaging subscores together into a single score per person)
- Computing the covariances of scores from one half with scores from the other half for every split
- Computing the variances of scores within each half for every split
- Computing the average split-half correlation with the average variances and covariance across all splits, using `corStatsByColumns()`

- Applying the Spearman-Brown formula to the absolute correlation using spearmanBrown(), and restoring the original sign after

cormean() was used to aggregate correlations in previous versions of this package & in the associated manuscript, but the method based on (co)variance averaging was found to be more accurate. This was suggested by prof. John Christie of Dalhousie University.

Value

A list containing

- r: the averaged reliability.
- ci: the 95% confidence intervals.
- allcors: a vector with the reliability of each iteration.
- nobs: the number of participants.
- scores: the individual participants scores in each split-half, contained in a list with two matrices (Only included if requested with include.scores).

Note

- This function can use a lot of memory in one go. If you are computing the reliability of a large dataset or you have little RAM, it may pay off to use the sequential version of this function instead: rapidsplit.chunks()
- It is currently unclear it is better to pre-process your data before or after splitting it. If you are computing the IAT D-score, you can therefore use errorhandling and standardize to perform these two actions after splitting, or you can process your data before splitting and forgo these two options.

Author(s)

Sercan Kahveci

References

Kahveci, S., Bathke, A.C. & Blechert, J. (2024) Reaction-time task reliability is more accurately computed with permutation-based split-half correlations than with Cronbach's alpha. Psychonomic Bulletin and Review. doi:10.3758/s1342302402597y

Examples

```
data(foodAAT)
# Reliability of the double difference score:
# [RT(push food)-RT(pull food)] - [RT(push object)-RT(pull object)]

frel<-rapidsplit(data=foodAAT,
                 subjvar="subjectid",
                 diffvars=c("is_pull","is_target"),
                 stratvars="stimid",
                 aggvar="RT",
                 splits=100)

print(frel)

plot(frel,type="all")
```

```
# Compute a single random split-half reliability of the error rate
rapidsplit(data=foodAAT,
           subjvar="subjectid",
           aggvar="error",
           splits=1,
           aggfunc="means")

# Compute the reliability of an IAT D-score
data(raceIAT)
rapidsplit(data=raceIAT,
           subjvar="session_id",
           diffvars="congruent",
           subscorevar="blocktype",
           aggvar="latency",
           errorhandling=list(type="fixedpenalty",errorvar="error",
                               fixedpenalty=600,blockvar="block_number"),
           splits=100,
           standardize=TRUE)


# Unstratified reliability of the median RT
rapidsplit.chunks(data=foodAAT,
                  subjvar="subjectid",
                  aggvar="RT",
                  splits=100,
                  aggfunc="medians",
                  chunks=8)

# Compute the reliability of Tukey's trimean of the RT
# on 2 CPU cores
trimean<-function(x){
  sum(quantile(x,c(.25,.5,.75))*c(1,2,1))/4
}
rapidsplit.chunks(data=foodAAT,
                  subjvar="subjectid",
                  aggvar="RT",
                  splits=200,
                  aggfunc=trimean,
                  cluster=2)
```

---

rapidsplithalf   *rapidsplithalf package*

Description

To learn more about rapidsplithalf, view the introductory vignette: `vignette("rapidsplithalf",package="rapidspl`

Author(s)

Sercan Kahveci

spearmanBrown	*Spearman-Brown correction Perform a Spearman-Brown correction on the provided correlation score.*

Description

Spearman-Brown correction Perform a Spearman-Brown correction on the provided correlation score.

Usage

```
spearmanBrown(r, ntests = 2, fix.negative = c("mirror", "nullify", "none"))
```

Arguments

r	To-be-corrected correlation coefficient.
ntests	An integer indicating how many times larger the full test is, for which the corrected correlation coefficient is being computed.
fix.negative	How will negative input values be dealt with?

- `"mirror"` submits the absolute correlations to the formula and restores the original sign afterwards.
- `"nullify"` sets negative correlations to zero.
- `"none"` leaves them as-is (not recommended).

Details

When `ntests=2`, the formula will compute what the correlation coefficient would be if the test were twice as long.

Value

Spearman-Brown corrected correlation coefficients.

Author(s)

Sercan Kahveci

Examples

```
spearmanBrown(.5)
```

---

stratifiedItersplits   *stratifiedItersplits*

Description

Generate stratified splits for a single participant

Usage

```
stratifiedItersplits(splits, groupsizes)
```

Arguments

splits   Number of iterations.

groupsizes   An integer vector of how many RTs per group need to be stratified.

Details

This equally splits what can be equally split within groups. Then it randomly splits all the leftovers to ensure near-equal split sizes. This function is moreso used internally, but you can use it if you know what you are doing.

Value

A matrix with zeroes and ones. Each column is a random split.

Examples

```
# We will create splits stratified by stimulus for a single participant
data(foodAAT)
currdata<-foodAAT[foodAAT$subjectid==3,]
currdata$stratfactor<-interaction(currdata$is_pull,currdata$is_target,currdata$stimid)
currdata<-currdata[order(currdata$stratfactor),]
groupsizes<-rle(as.character(currdata$stratfactor))$lengths

mysplits<-stratifiedItersplits(splits=1000,groupsizes=groupsizes)

# Now the data can be split with the values from any column.
half1<-currdata[mysplits[,1]==1,]
half2<-currdata[mysplits[,1]==0,]

# Or the split objects can be used as masks for the aggregation functions in this package
meansByMask(x=currdata$RT,mask=mysplits==1)
```

Index